

# ZIL Performance: How I Doubled Sync Write Speed

# Agenda

1. What is the ZIL?
2. How is it used? How does it work?
3. The problem to be fixed; the solution.
4. Details on the changes I made.
5. Performance testing and results.

# 1 – What is the ZIL?

# What is the ZIL?

- ZIL: Acronym for (Z)FS (I)ntent (L)og
  - Logs synchronous operations to disk, before `spa_sync()`
  - What operations get logged?
    - `zfs_create`, `zfs_remove`, `zfs_write`, etc.
    - Doesn't include non-modifying ZPL operations:
      - `zfs_read`, `zfs_seek`, etc.
  - What gets logged?
    - The fact that a logical operation is occurring is logged
      - `zfs_remove` → directory object ID + name only
    - Not logging which blocks will change due to logical operation

# When is the ZIL used?

- Always<sup>\*</sup>
  - ZPL operations (itx's) logged via in-memory lists
  - lists of in-memory itx's written to disk via `zil_commit()`
  - `zil_commit()` called for:
    - *any* sync write<sup>\*\*</sup>

<sup>\*</sup>Except when dataset configured with: `sync=disabled`. <sup>\*\*</sup>Except when dataset configured with: `sync=always`.

# What is the SLOG?

- SLOG: Acronym for (S)eparate (LOG) Device
  - An SLOG is not necessary
    - By default (no SLOG), ZIL will write to main pool VDEVs
  - An SLOG can be used to improve latency of ZIL writes
    - When attached, ZIL writes to SLOG instead of main pool\*
- Conceptually, SLOG is different than the ZIL
  - ZIL is mechanism for writing, SLOG is device written to
- ZIL is used, even if no SLOG attached


\*For some operations; see code for details.

# Why does the ZIL exist?

- Writes in ZFS are "write-back"
  - Data is first written and stored in-memory, in DMU layer
  - Later, data for whole pool written to disk via `spa_sync()`
- Without the ZIL, sync operations could wait for `spa_sync()`
  - `spa_sync()` can take tens of seconds (or more) to complete
- Further, with the ZIL, write amplification can be mitigated
  - A single ZPL operation can cause many writes to occur
  - ZIL allows operation to "complete" with minimal data written
- ZIL needed to provide "fast" synchronous semantics to applications
  - Correctness could be achieved without it, but would be "too slow"

# ZIL On-Disk Format

- Each dataset has it's own unique ZIL on-disk
- ZIL stored on-disk as a singly linked list of ZIL blocks (lwb's)





## 2 – How is the ZIL used?

# How is the ZIL used?

- ZPL will generally interact with the ZIL in two phases:
  1. Log the operation(s) — `zil_itx_assign`
    - Tells the ZIL an operation is occurring
  2. Commit the operation(s) — `zil_commit`
    - Causes the ZIL to write log record of operation to disk

# Example: zfs\_write

- `zfs_write` → `zfs_log_write`
- `zfs_log_write`
  - `zil_itx_create`
  - `zil_itx_assign`
- `zfs_write` → `zil_commit`
- Most ZPL operations have a corresponding `zfs_log_*` function
  - `zfs_log_create`
  - `zfs_log_remove`
  - `zfs_log_link`
  - `zfs_log_symlink`
  - `zfs_log_truncate`
  - `zfs_log_setattr`
  - ...

# Example: zfs\_fsync

- `zfs_fsync` → `zil_commit`
  - `fsync` doesn't create any new modifications
  - only writes previous itx's to disk
    - thus, no `zfs_log_fsync` function

# Contract between ZIL and ZPL.

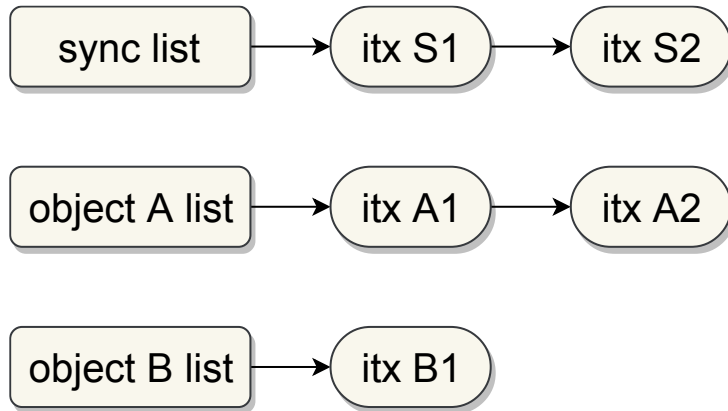
- Parameters to `zil_commit`: ZIL pointer, object number
  - These uniquely identify an object whose data is to be committed
- When `zil_commit` returns:
  - Operations *relevant* to the object specified, will be *persistent* on disk
  - relevant – all operations that would modify that object
  - persistent – Log block(s) written (completed) → disk flushed
- Interface of `zil_commit` doesn't specify *which* operation(s) to commit

## 2 – How does the ZIL work?

# How does the ZIL work?

- In memory ZIL contains per-txg `itxg_t` structures
- Each `itxg_t` contains:
  - A single list of sync operations (for all objects)
  - Object specific lists of async operations

# Example: itx lists





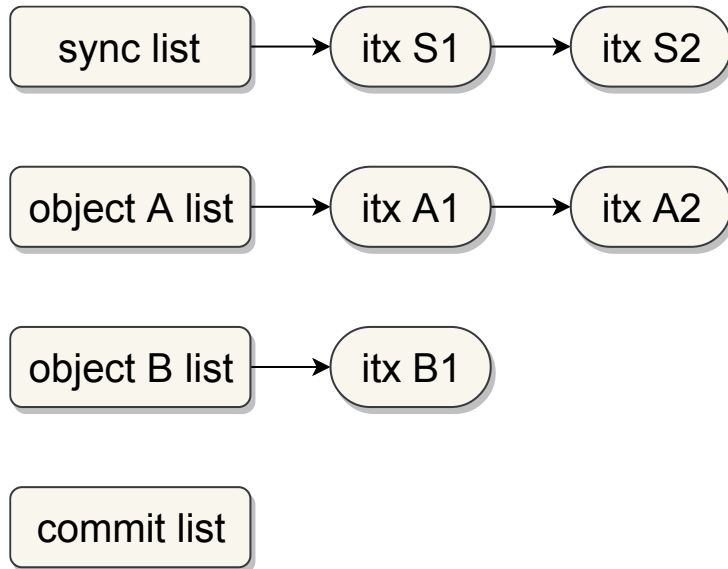
# How are itx's written to disk?

- `zil_commit` handles the process of writing `itx_t`'s to disk:

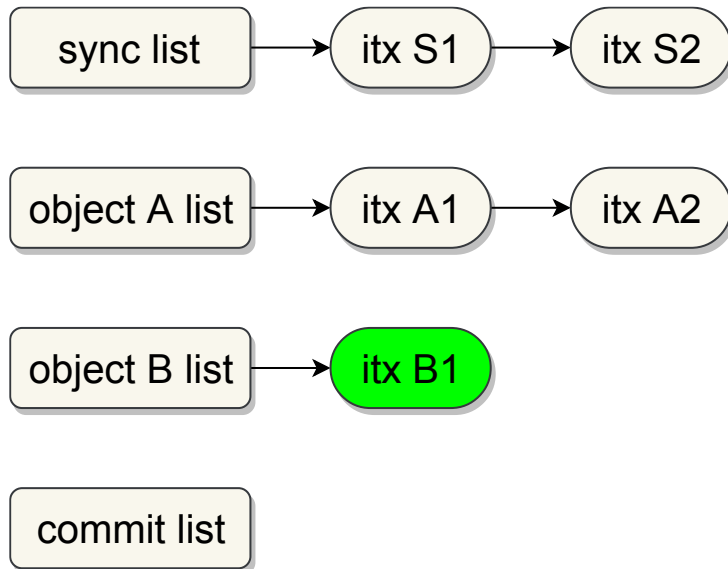
# How are itx's written to disk?

- `zil_commit` handles the process of writing `itx_t`'s to disk:
  1. find all relevant `itx`'s, move them to the "commit list"

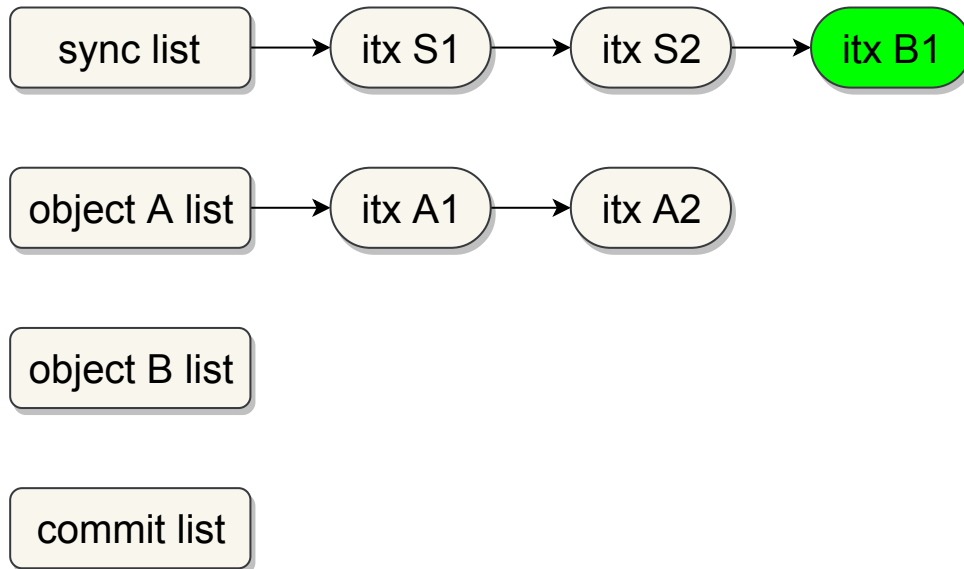
# Example: zil\_commit Object B



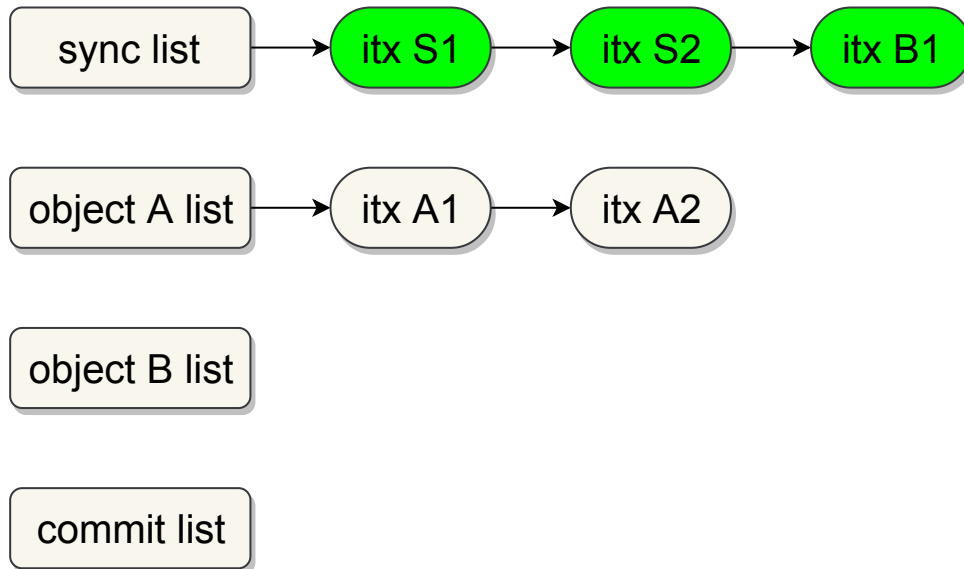
# Example: zil\_commit Object B



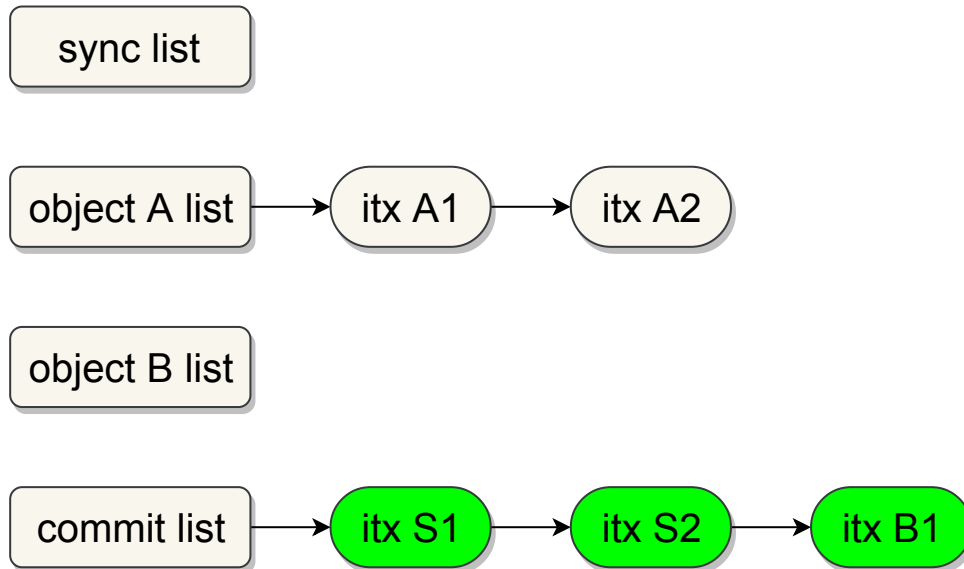
# Example: zil\_commit Object B



# Example: zil\_commit Object B



# Example: zil\_commit Object B

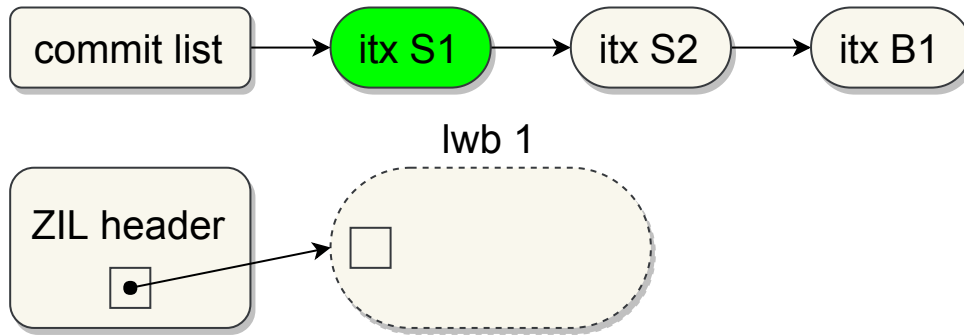


# How are itx's written to disk?

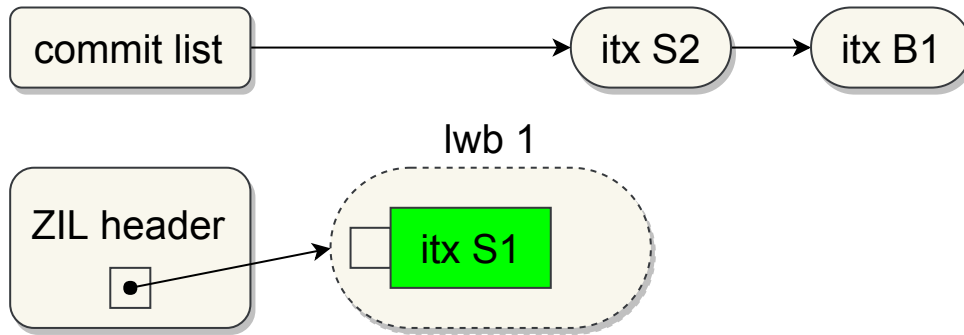
- `zil_commit` handles the process of writing `itx_t`'s to disk:
  1. ~~Move async itx's for object being committed, to the sync list~~
  2. Write all commit list itx's to disk




# Example: zil\_commit Object B



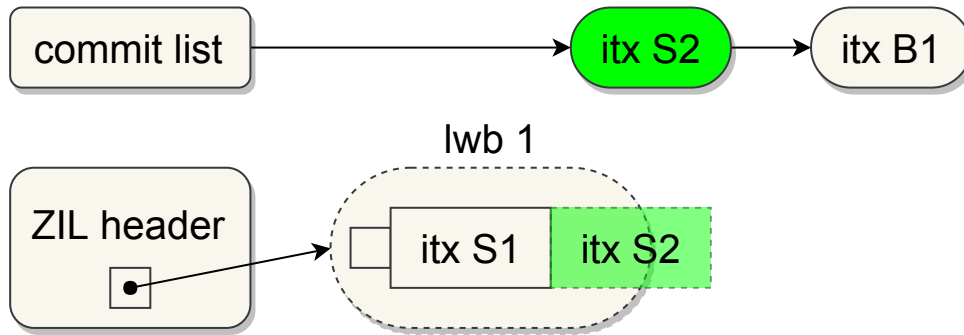
# Example: zil\_commit Object B



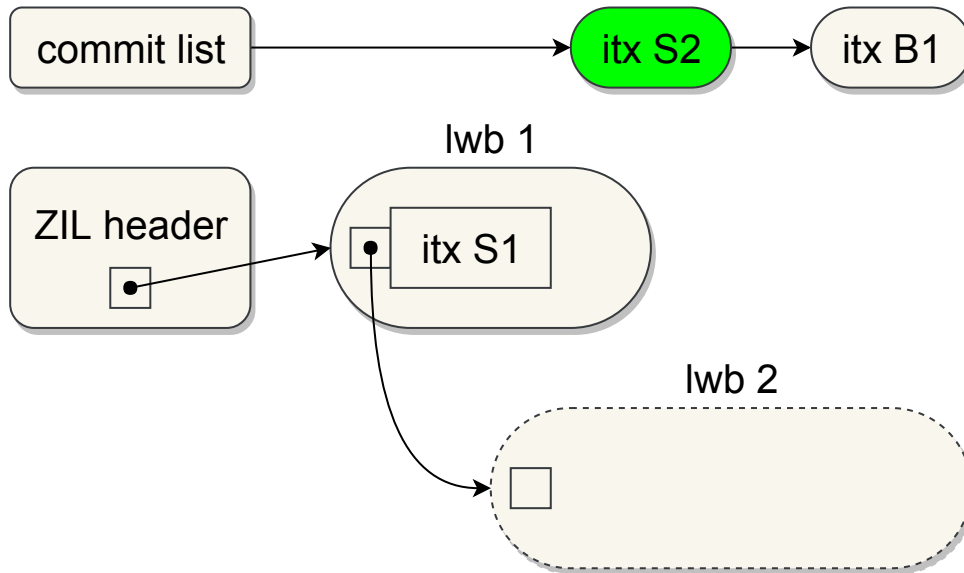
# Example: zil\_commit Object B



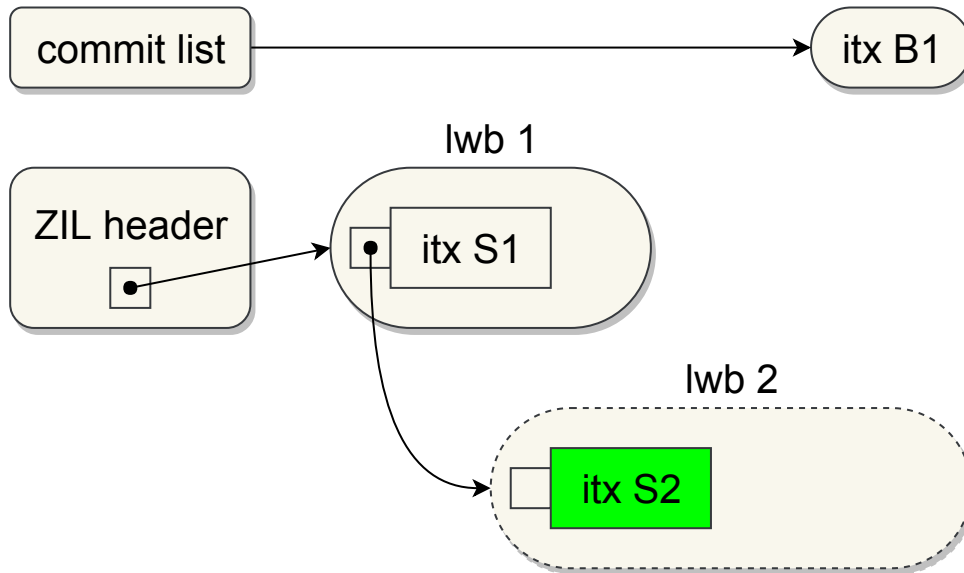
# Example: zil\_commit Object B



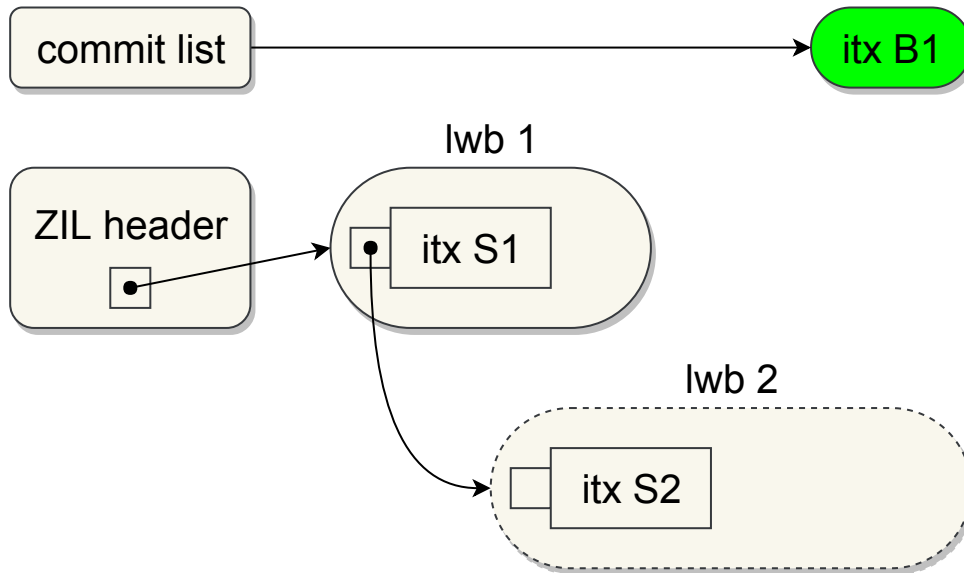
# Example: zil\_commit Object B



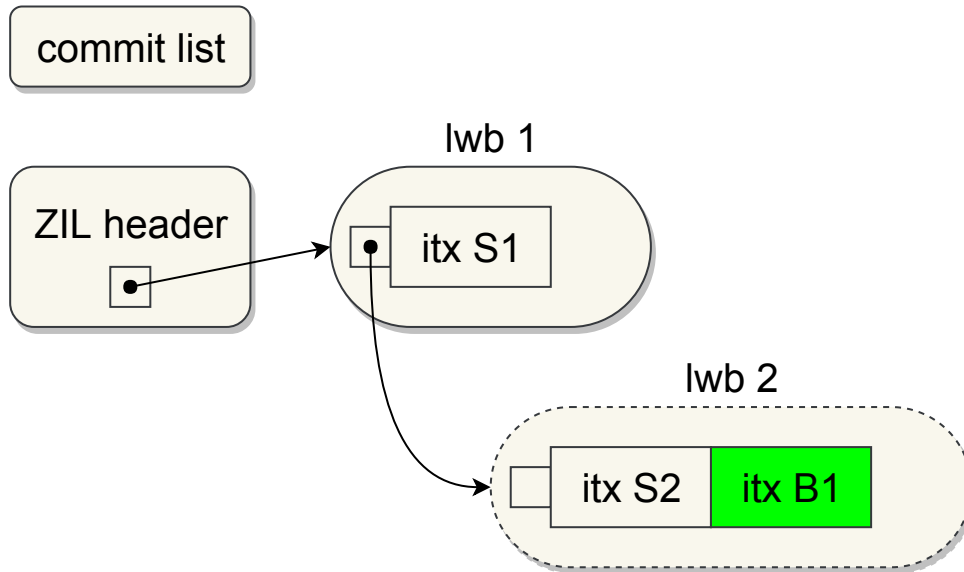
# Example: zil\_commit Object B



# Example: zil\_commit Object B

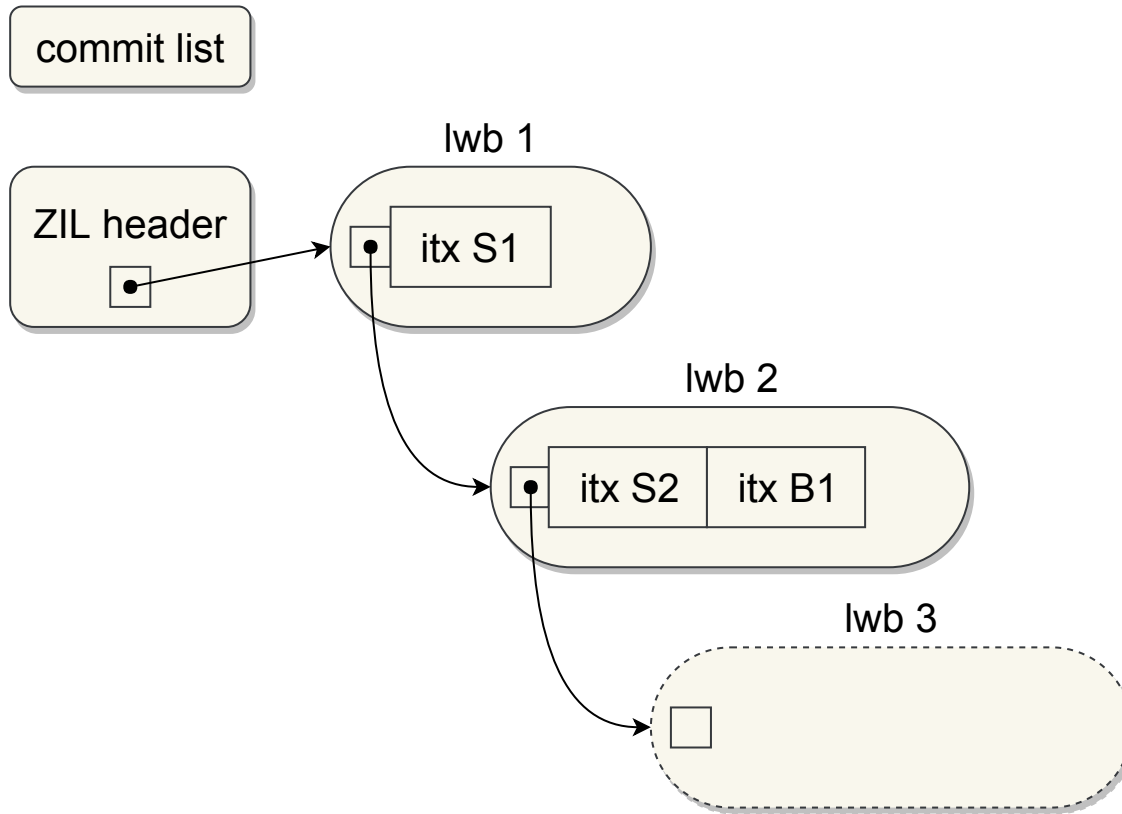


# Example: zil\_commit Object B





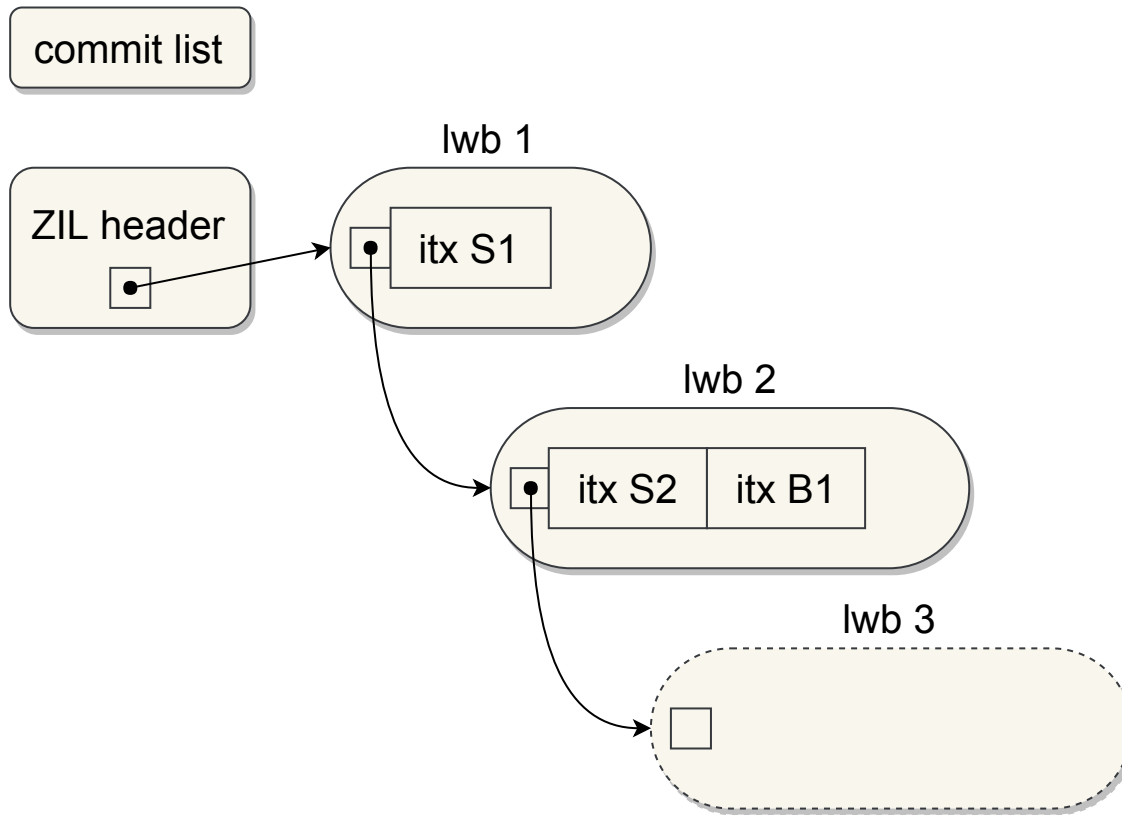
# Example: zil\_commit Object B



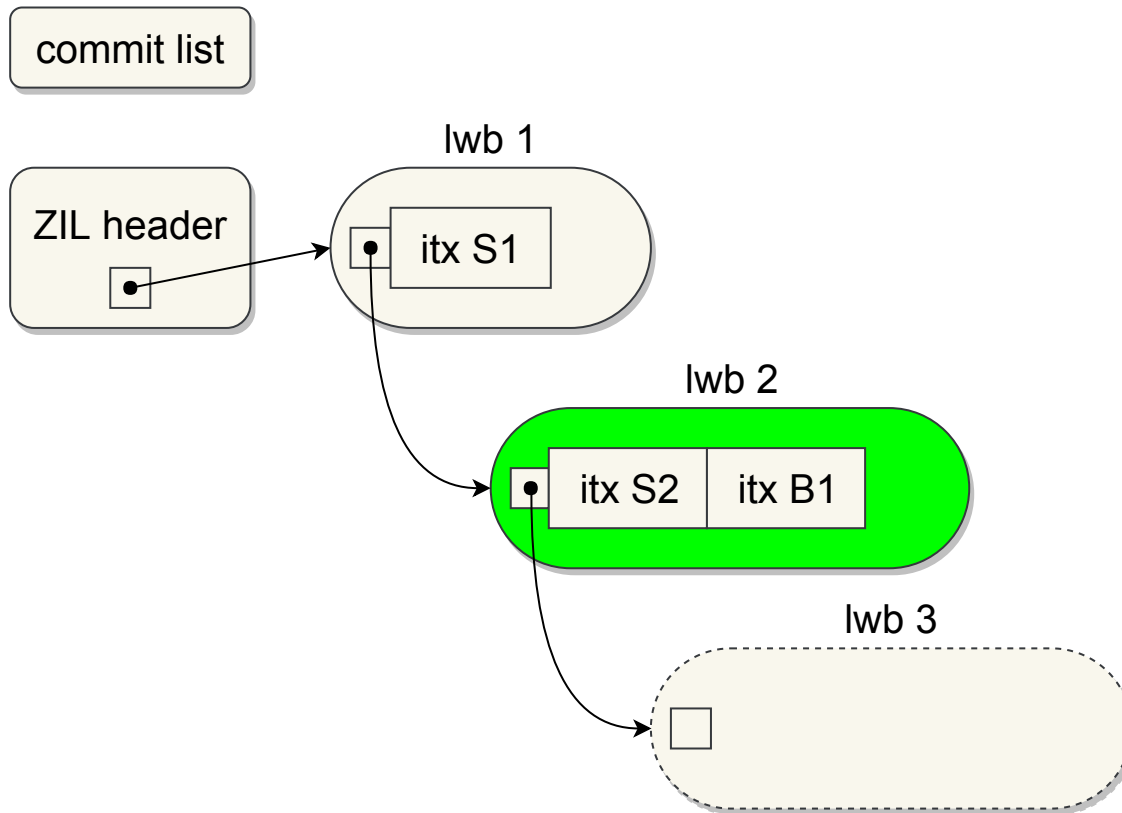
# How are itx's written to disk?

- `ztl_commit` handles the process of writing `itx_t`'s to disk:
  1. ~~Move async itx's for object being committed, to the sync list~~
  2. ~~Write all commit list itx's to disk~~
  3. Wait for all ZIL block writes to complete

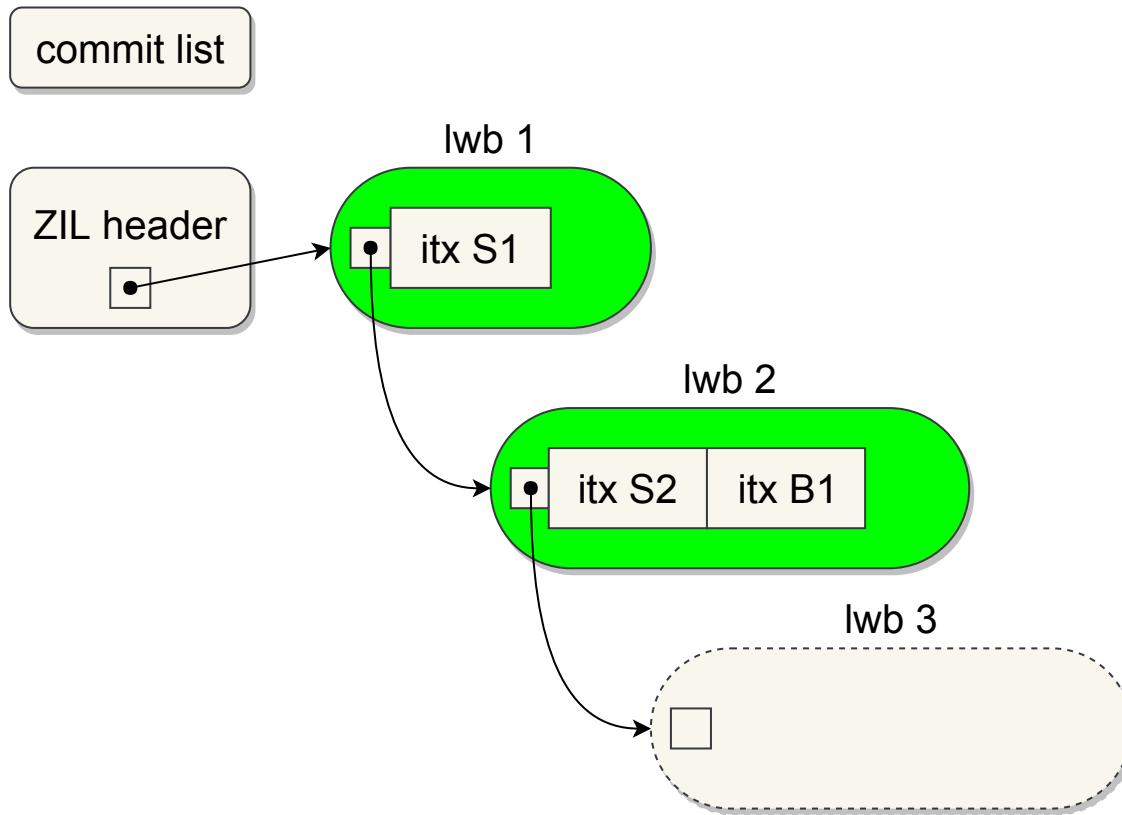
# Example: zil\_commit Object B



# Example: zil\_commit Object B



# Example: zil\_commit Object B



# How are itx's written to disk?

- `zil_commit` handles the process of writing `itx_t`'s to disk:
  1. ~~Move async itx's for object being committed, to the sync list~~
  2. ~~Write all commit list itx's to disk~~
  3. ~~Wait for all ZIL block writes to complete~~
  4. Flush VDEVs

# How are itx's written to disk?

- `zil_commit` handles the process of writing `itx_t`'s to disk:
  1. ~~Move async itx's for object being committed, to the sync list~~
  2. ~~Write all commit list itx's to disk~~
  3. ~~Wait for all ZIL block writes to complete~~
  4. ~~Flush VDEVs~~
  5. Notify waiting threads

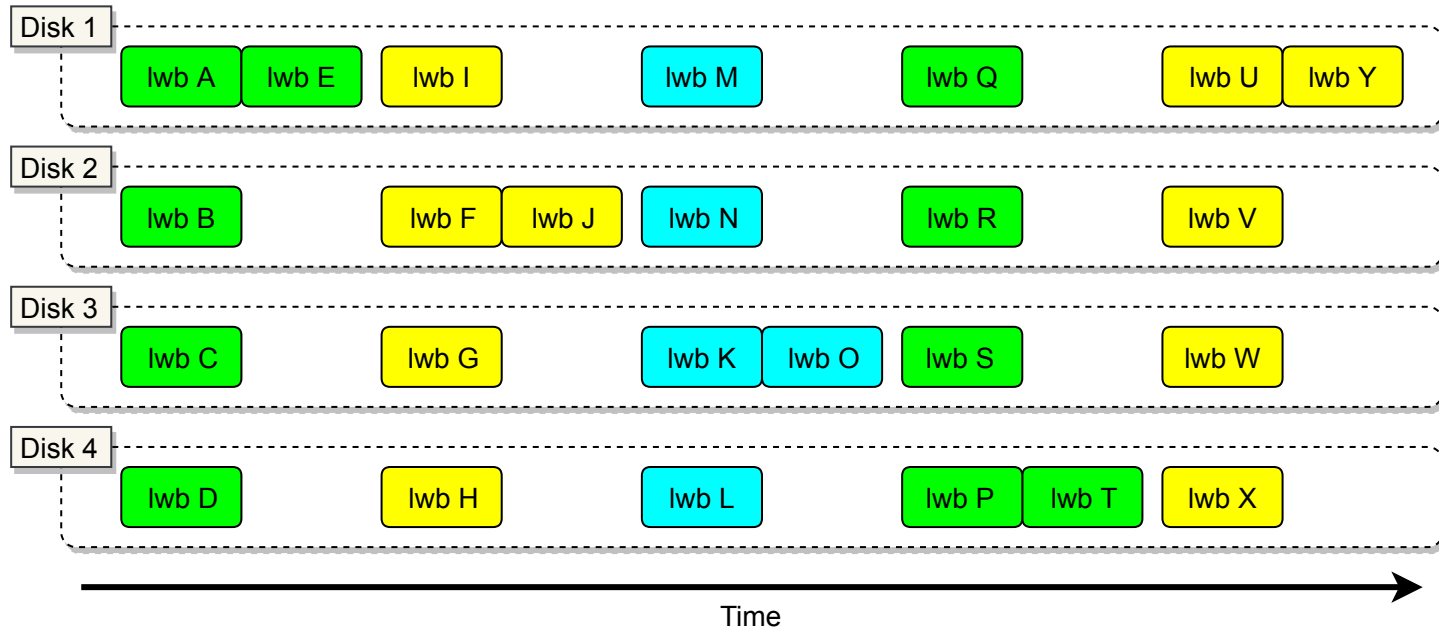
# 3 – Problem



# Problem

1. `itx`'s grouped and written in "batches"
  - The commit list constitutes a batch
  - Batch size proportional to sync workload on system
2. Waiting threads only notified when *all* ZIL blocks in batch complete
3. Only a single batch processed at a time

# Problem



- Time spent servicing lwb's for each disk
- Color indicates order waiting threads notified

# Implications


1. `zil_commit` latency proportional to system workload, *not* disk latency
  - Fast SLOG may not compensate for large workload
2. Disk "anomalies" → larger batches → increased `zil_commit` latency
  - e.g. temporary network delays when using network storage
3. New calls to `zil_commit` wait for "current" batch, *and* "next" batch
  - Average `zil_commit` latency equal to latency of 1.5 batches

# 3 – Solution

# Solution


- Remove concept of "batches":
  1. Allow `ztl_commit` to issue new ZIL block writes immediately
    - In contrast to waiting for the current batch to complete
  2. Notify threads immediately when *dependent* lwb's on disk
    - In contrast to waiting for *all* lwb's on disk

# Problem



- Time spent servicing lwb's for each disk
- Color indicates order waiting threads notified

# Solution




- Time spent servicing lwb's for each disk
- Color indicates order waiting threads notified

## 4 – Details on the Changes I Made

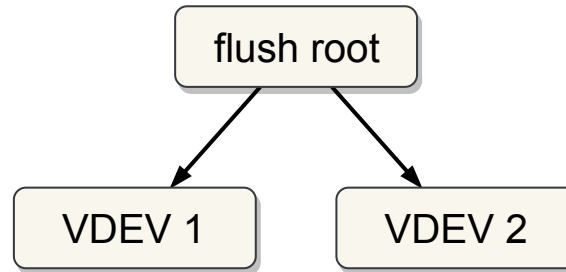


Before


Step 1



Step 2

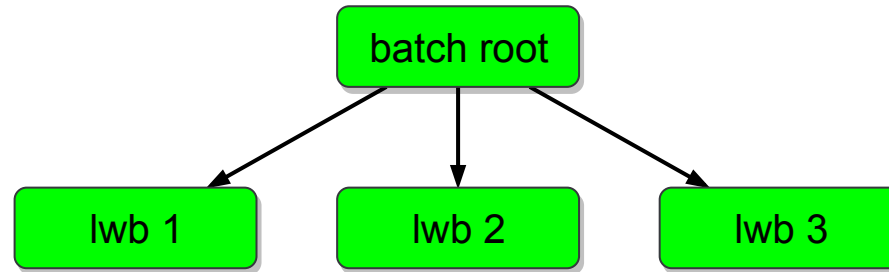


Step 3




Before


Step 1



Step 2

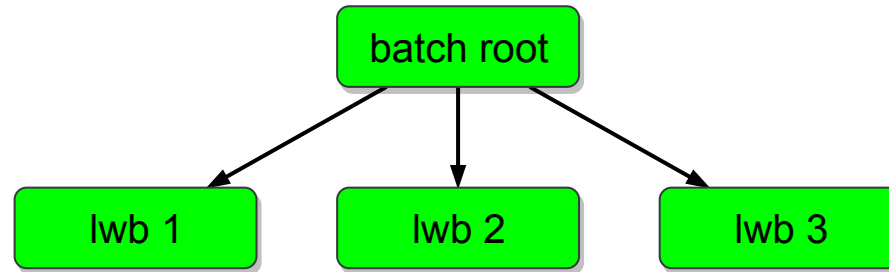


Step 3

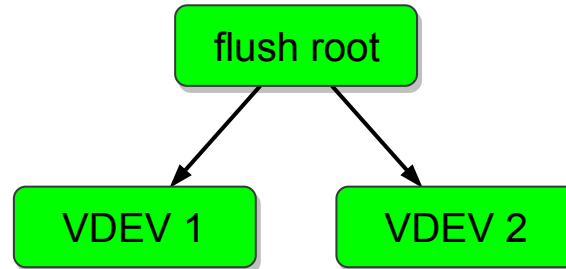


Before


Step 1



Step 2




Step 3




Before


Step 1



Step 2




Step 3




Before


Step 1



Step 2




Step 3




Before


Step 1



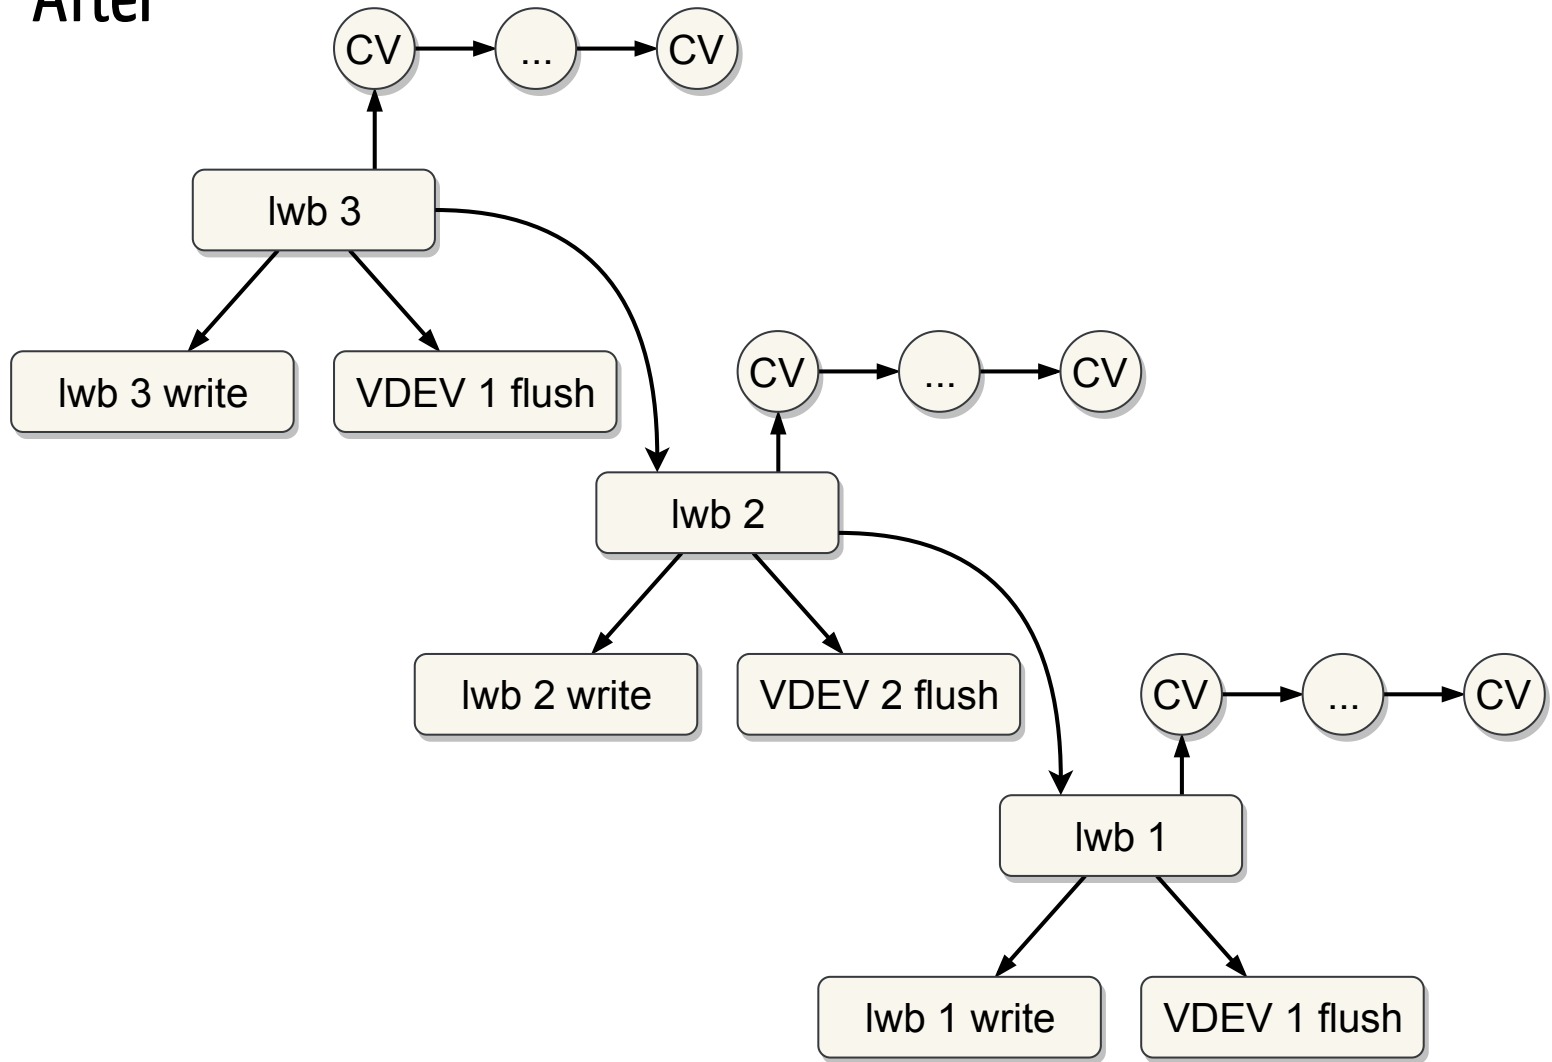
Step 2



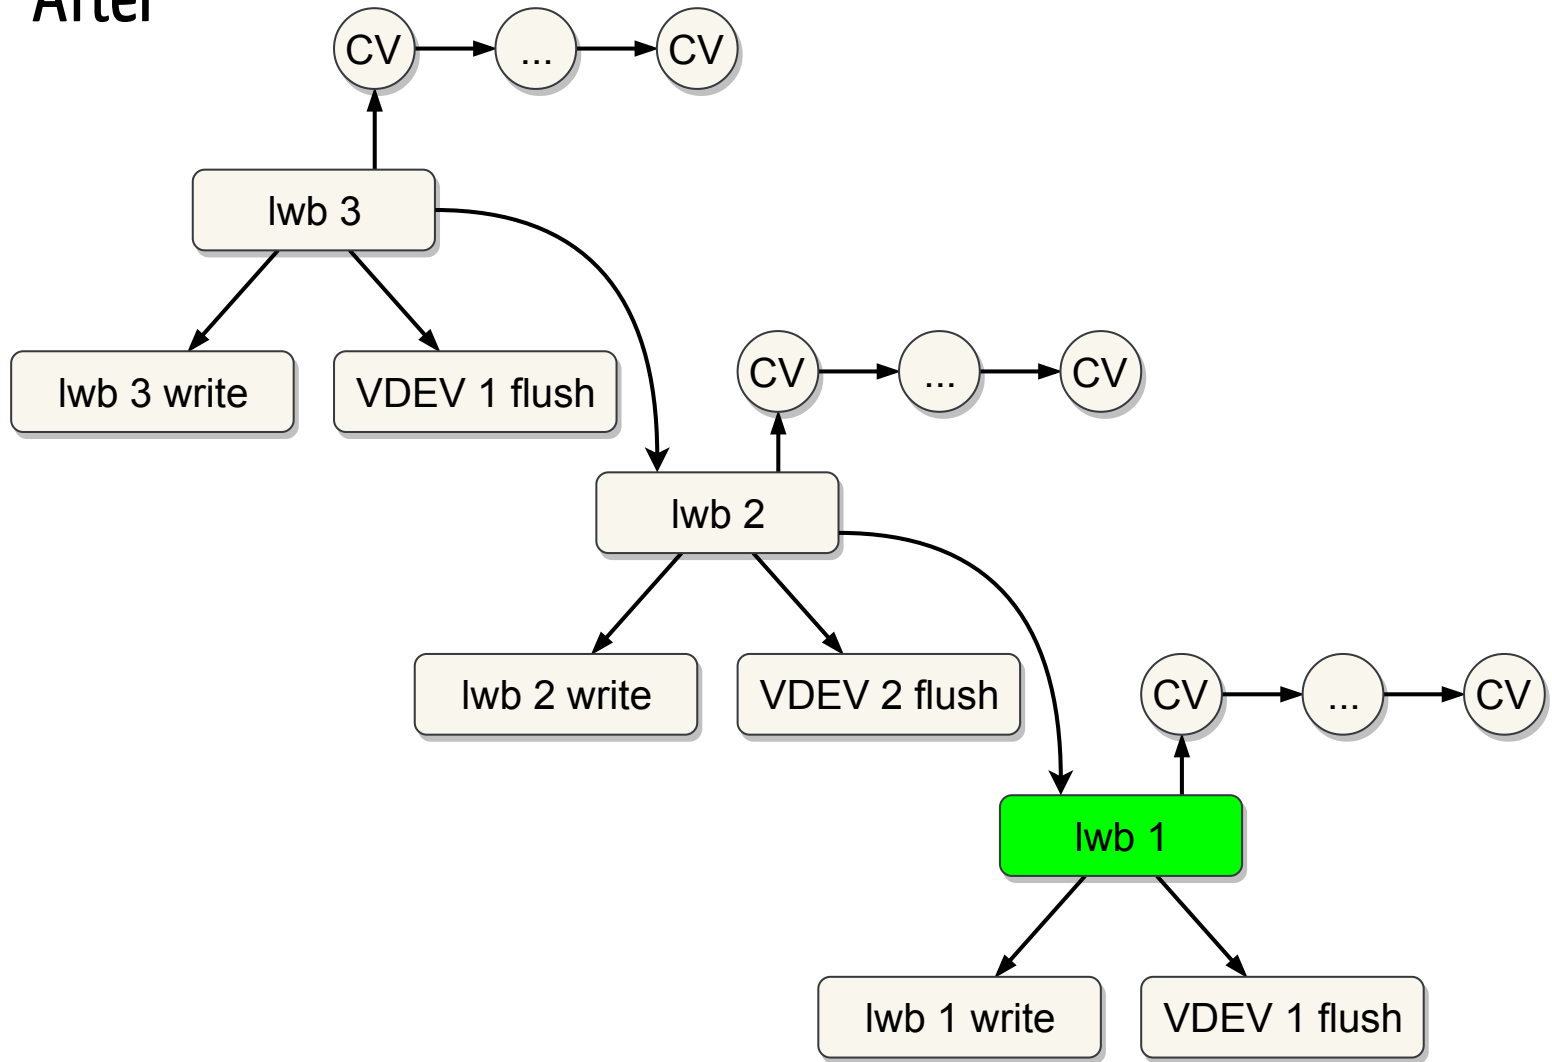
Step 3



After




After

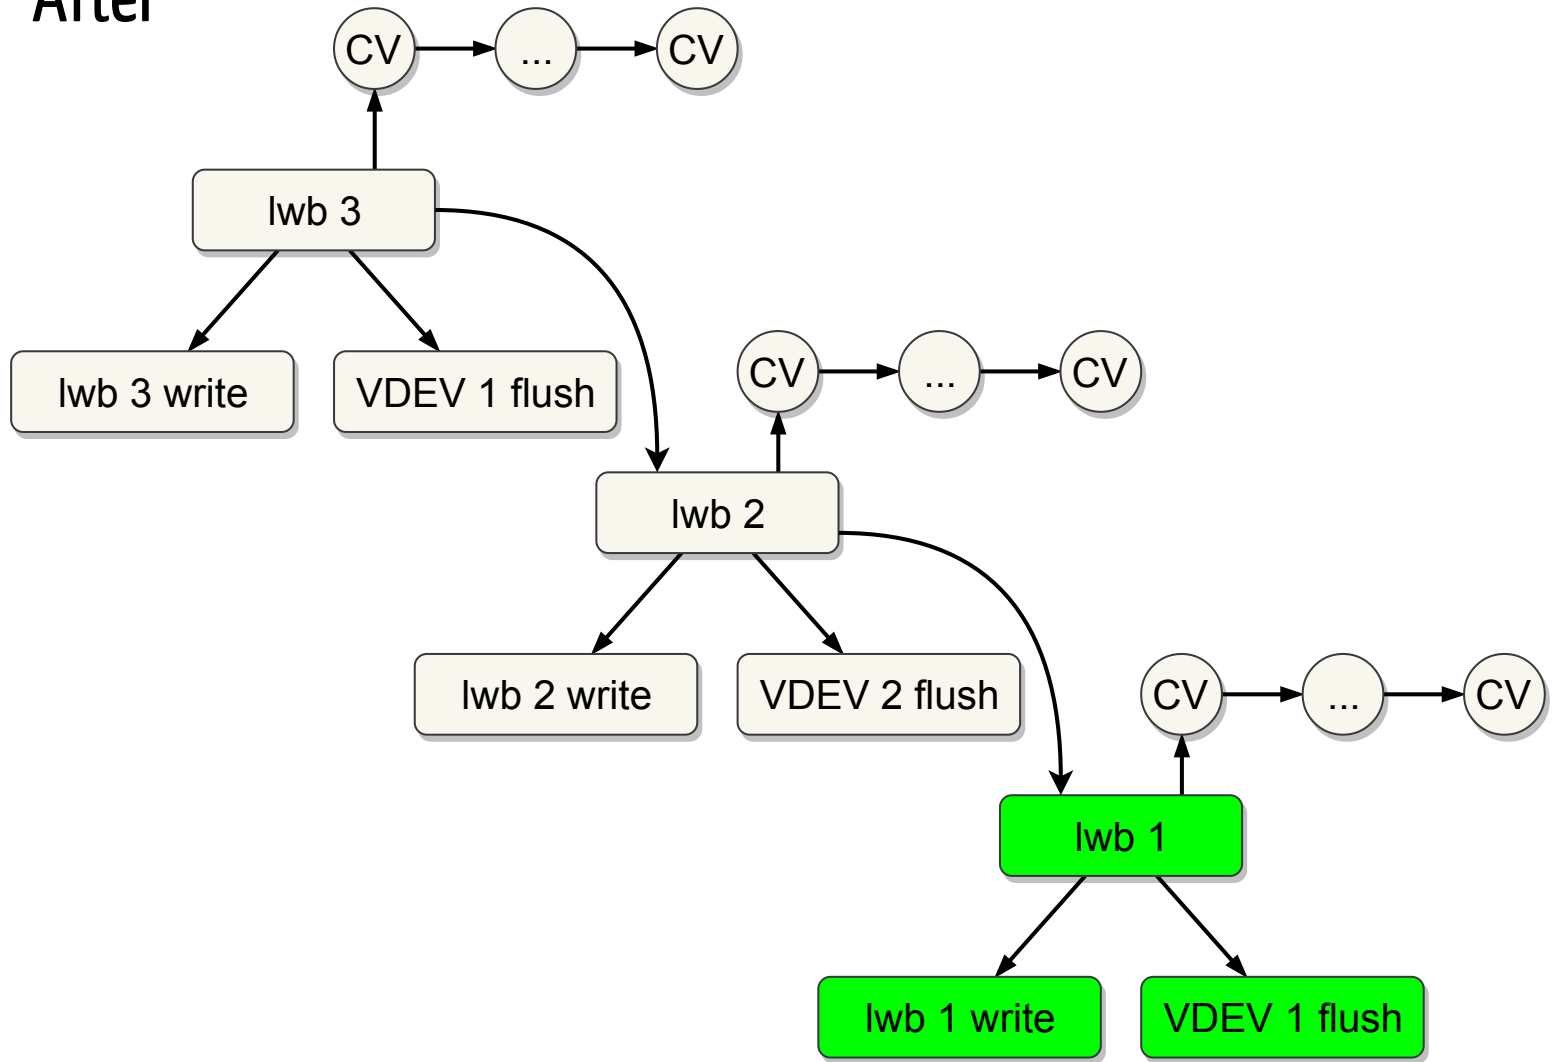




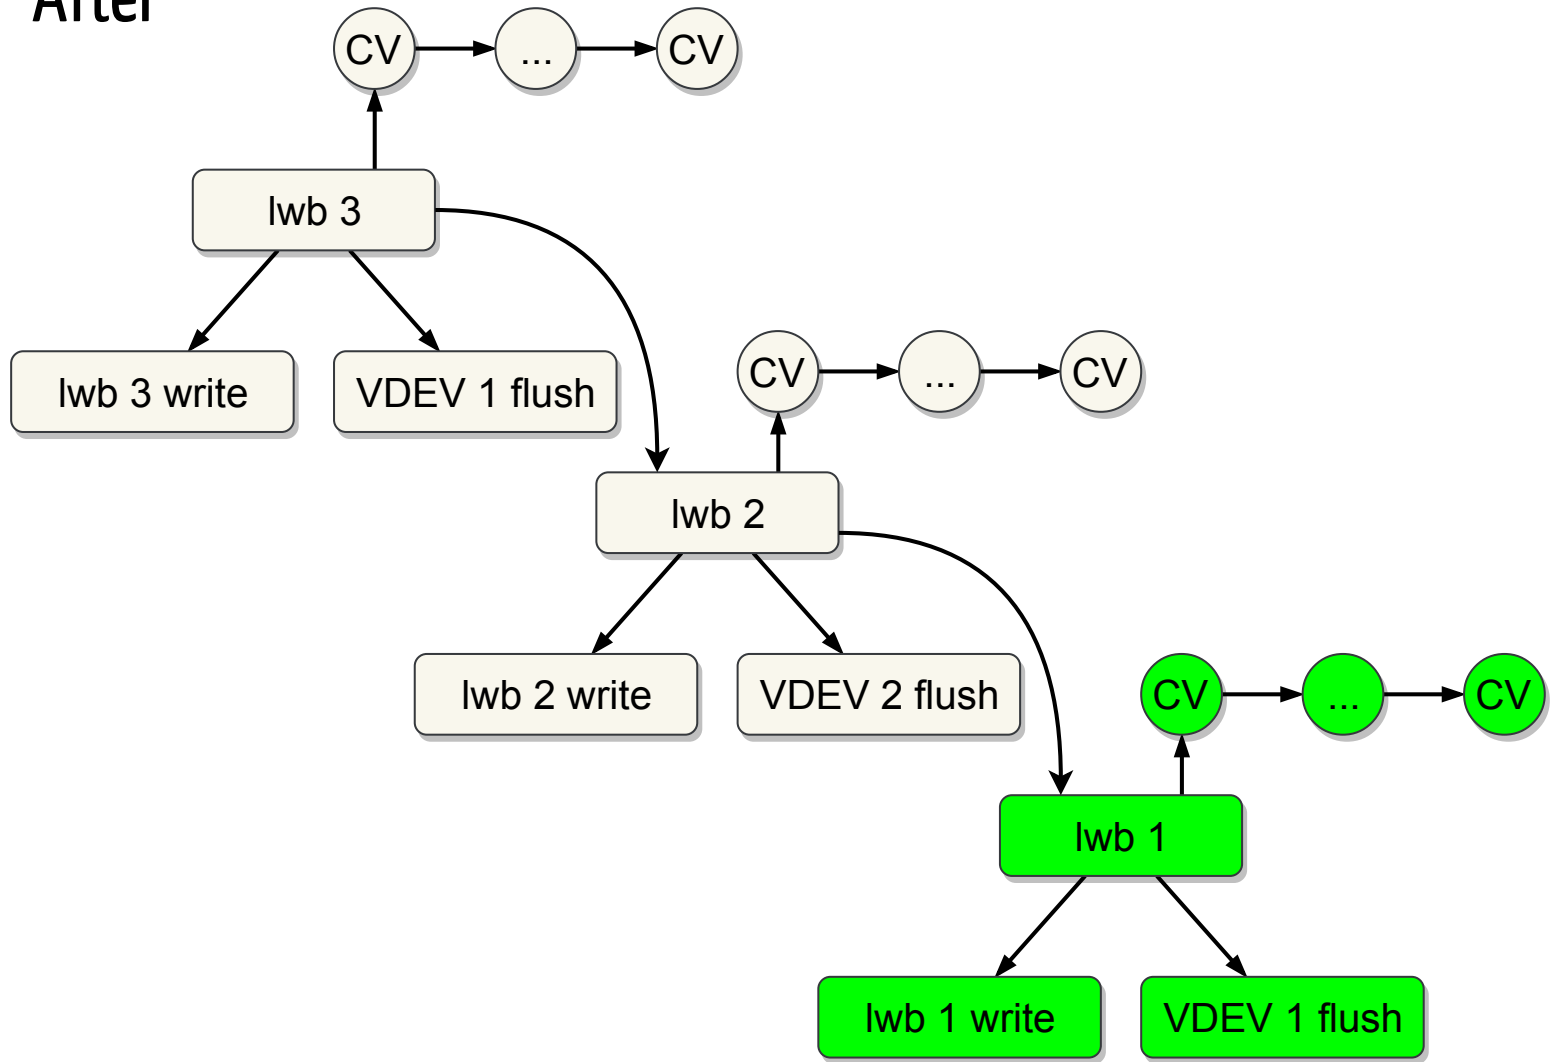
After




After




After



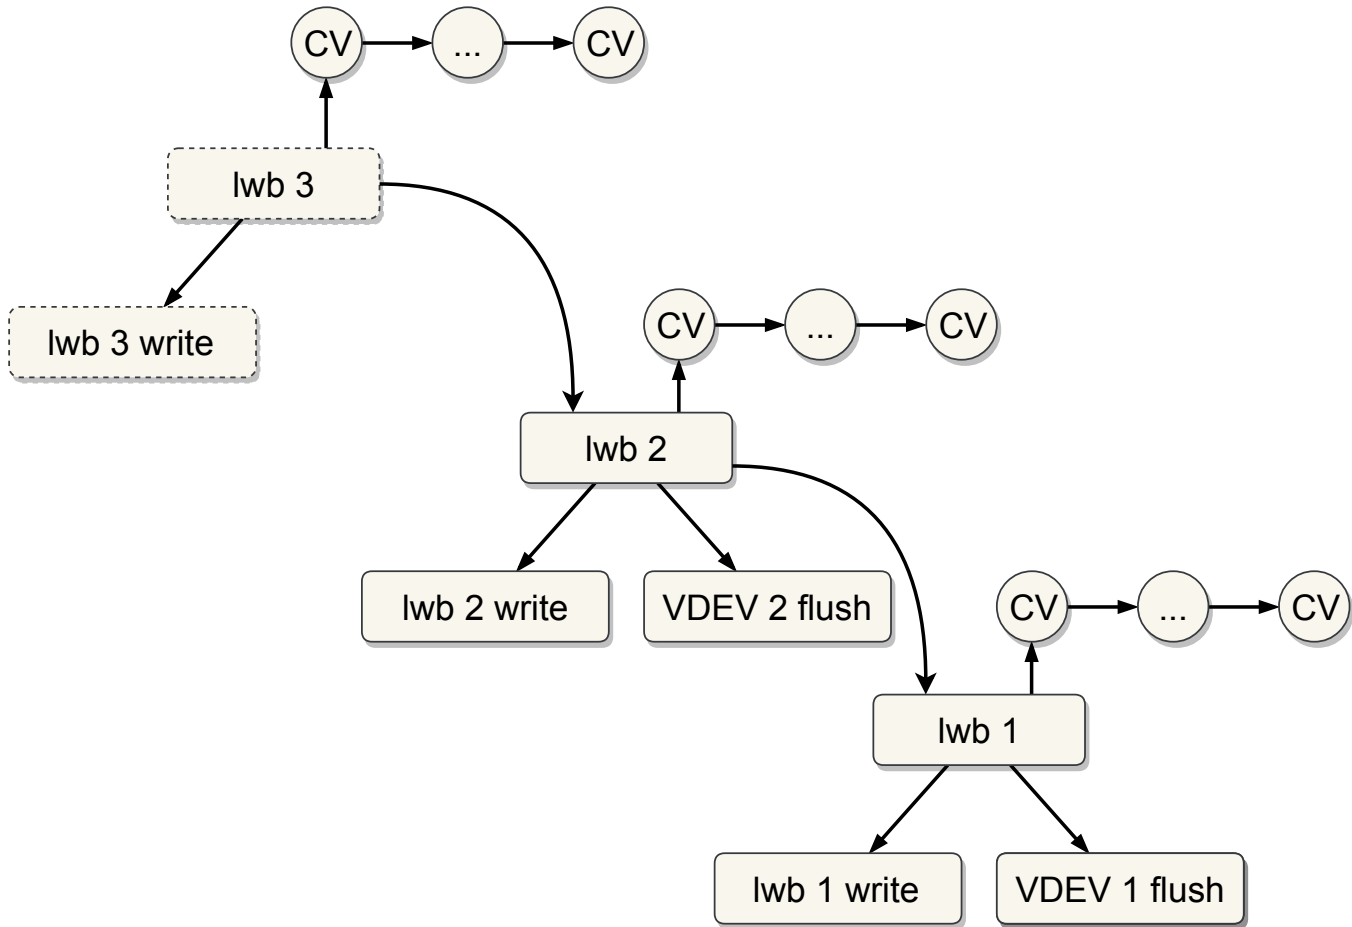
After



After




# New Tunable: lwb Timeout



\*New tunable named: `zfs_commit_timeout_pct`


# 5 – Performance testing and results

# ~83% Increase in IOPs on Average – Max Rate – 8 HDDs






# ~48% Increase in IOPs on Average – Max Rate – 8 SSDs




# ~27% Decrease in Latency on Average – Fixed Rate – 8 HDDs



\*IOPs increased with new code, and >64 threads; those data points omitted.

# ~16% Decrease in Latency on Average – Fixed Rate – 8 SSDs



# More Details

- Two fio workloads were used:
  1. each thread submitting sync writes as fast as it could
  2. each thread submitting 64 sync writes per second
- 1, 2, 4, and 8 disk zpools; both SSD and HDD
- fio threads ranging from 1 to 1024; increasing in powers of 2
- Full details can be found [here](#)

End